

| PROJECT 4: FILE SYSTEMS |

---- GROUP 12----

Sakaar Khurana 10627
Nikhil Aggarwal 10446
Aditi Krishn 10039

INDEXED AND EXTENSIBLE FILES =====

---- DATA STRUCTURES ----

1. in inode.c

1.1 #define INODE_BLOCKS 124

This is the number of direct blocks in an on disk inode

1.2

```
int INODE_LIST = BLOCK_SECTOR_SIZE / sizeof(block_sector_t);
```

This is the number of blocks of the indirect inode blocks

1.3

```
struct inode_disk
{
    unsigned number_blocks;
    off_t length;                /* File size in bytes. */
    unsigned magic;              /* Magic number. */
    block_sector_t sector[INODE_BLOCKS]; /* for the block numbers */
    enum file_type type;
};
```

The new inode_disk structure which contains simple information about the file it represents, length, numberofblocks it occupies, a sector array for direct, 1 indirect, and 1 double indirect block.

1.4

/* On-disk inode list for linked and or double linked lists.

Must be exactly BLOCK_SECTOR_SIZE bytes long. */

```
struct inode_list
```

```
{
    block_sector_t sector[128];    /* for the block numbers */
};
```

This represents the linked or double linked block that I store on the disk.

1.5 in struct inode:

```
//struct inode_disk data;          /* Inode content. */
struct lock inode_lock;
```

We removed the reference to inode disk in the inode because it will be accessed over the cache and the sector in the inode.

We added a lock to the structure for synchronization

The maximum size of a file is:

- A) direct blocks 124
- B) single indirect blocks 128
- C) double indirect blocks 128 * 128

=> $124 + 128 + 128 * 128 = 16636$ Blocks which is
 $16636 * 512 = 8.517.632$ bytes or 8.318 MB

---- Multi level indexing ----

We support multilevel index up to one doubly indirect block.
 It is the standard way in linux and easy to manage.
 The blocks can be allocated and extended in a consistent way.

---- TEST CASES ----

- All the files when are brought into the filesystem call the function `file_grow` in `filesys/inode.c` where data is allocated as per need.
 eg. run following from `userprog/build`:
`pintos -p ../../examples/echo -a echo -- -q`
- Files can also be rown after cretion time usig same general function.

SUBDIRECTORIES =====

---- DATA STRUCTURES ----

1. in `inode.h`

```
enum file_type
{
    FILE_FILE,
    FILE_DIR
};
```

This enumeration tells if an inode is a directory or file

2. in `thread.h`, structure `thread`:

```
struct dir * pwd;
```

The current working directory of the thread.

---- ALGORITHMS ----

- First of all we call a parser function to remove any multiple entries of `/` in the filename. (in `filesys.c`, `parse_filename()`).
- Then we call the function `dir_lookup_rec()` in `directory.c` which will does the traversing.
- It will look for a `/` in the first character of the filename, if it exists
 we know that we have an absolute path and start at root otherwise use the current working directory as starting point.
- We strtok the filename and walk through the directories in the tokens one by one to see if they are a directory.

---- Explaining the Directory Structure ----

- We are saving a struct dir in the thread structure that represent the current working directory.
- This way it is easy for us to use it when a user opens a file or directory with relative pathname.
- We just grab the directory and go from there.

---- TEST CASES ----

- mkdir.c: makes a directory passed to it.
- chdir.c: changes the directory to the directory passed to it.
- open.c: changes to a directory 'dir', makes 2 directories over there, changes directory back to the root directory, opens a file with directory 'dir' and calls readdir over there.
This test case tests the functionality of all the system-calls.

BUFFER CACHE =====

---- DATA STRUCTURES ----

1. in filesystem/cache.c

1.1

```
struct cache_block {
    //block sector on disk
    block_sector_t bid;
    //corresponding kernel page
    void * kpage;
    //fields to check access and if someone wrote to the page
    bool dirty;
    bool accessed;
    int reader;
    int writer;
};
```

This represents an entry in the cache table.
It has the corresponding block number on the disk, a kpage for the kernel page where it is stored right now.
A dirty and accessed variable for writing and eviction and reader and writer are used for synchronization.

1.2

```
//how many blocks is one page
int SECPP;
```

The global variable tells the number of blocks that fit in one kernel page.

1.3

```
//cache array
struct cache_block * cache[cache_size];
```

This is the actual cache table.

1.4

```
//bitmap to identify free entries
struct bitmap * cache_table;
```

The bitmap that is used to identify empty entries in the cache table.

1.5

```
struct lock cache_globallock;
```

A global lock to protect the cache table when I add new elements or

2. in filesystem/inode.h

2.1

```
//for the read ahead thread
struct lock lock_readahead;
```

A lock for the read ahead thread.

2.2

```
struct list list_readahead;
```

A list for the read ahead thread.

2.3

```
struct condition cond_readahead;
```

A condition variable to signal the waiting read ahead thread to wake up.

2.4

```
//structure for the readahead list
struct readahead {
    block_sector_t bid;
    struct list_elem elem;
};
```

An element in the read ahead list.

---- ALGORITHMS ----

- Eviction Algorithm:

- We use the second chance algorithm from project 3. Every block has an accessed variable. A prerequisite for a page in order to get evicted is that the page has no reader or writer at the moment.
- If the accessed bit is true We reset it to fault and leave the block in the list.
- If one block has a false accessed bit we evict it.

---- Working of Buffer cache ----

- When a process needs to access the same part of a file multiple times or when multiple processes access the same part of a file a lot then caching will be a lot faster than reading or writing to the disk directly.
- Also when a process changes a lot of bytes in the same block a cache will be useful to prevent slow reads and writes from the disk.

- Read ahead will have its benefits when a process needs to read contiguous blocks from a file, so from 1 to last byte.

---- TEST CASES ----

Interaction between kernel and filesystem occurs through buffer cache, so all user programs running properly signifies correct working of buffer cache.